

```
In [122... # import numerical libraries
import numpy as np
import pandas as pd
import warnings

# import graphical plot libraries
import seaborn as sn
import matplotlib.pyplot as plt
%matplotlib inline

# Import Linear regression machine Learning Libraries
from sklearn import preprocessing
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score
```

```
In [124... warnings.filterwarnings('ignore')
# read / Load data from csv.
data = pd.read_csv(r"C:\sample\datafiles\car-mpg.csv")
data.head()
```

```
Out[124...      mpg  cyl  disp  hp  wt  acc  yr  origin  car_type      car_name
0   18.0    8  307.0  130 3504  12.0  70     1     0  chevrolet chevelle malibu
1   15.0    8  350.0  165 3693  11.5  70     1     0    buick skylark 320
2   18.0    8  318.0  150 3436  11.0  70     1     0  plymouth satellite
3   16.0    8  304.0  150 3433  12.0  70     1     0    amc rebel sst
4   17.0    8  302.0  140 3449  10.5  70     1     0    ford torino
```

```
In [126... #Drop car name
#replace origin 1,2,3...dont forget get dummies
#Replace ? with nan
#Replace all nan with median

#data = data.drop(['car_name'], axis = 1)
#data['origin'] = data['origin'].replace({1:'america', 2:'europe', 3:'asia'})
#data = pd.get_dummies(data,columns = ['origin'])
#data = data.replace('?', np.nan)
#data = data.apply(lambda x: x.fillna(x.median()), axis = 0)

data = data.drop(['car_name'], axis = 1)
data['origin'] = data['origin'].replace({1: 'america', 2: 'europe', 3: 'asia'})
data = pd.get_dummies(data,columns = ['origin'])
data = data.replace('?', np.nan)
# Identify columns that should be numeric but might contain NaNs or be object type.
# It's good practice to list your numeric columns or infer them.
# For example, if 'horsepower' and 'acceleration' were originally numbers but had
numeric_cols = ['cyl', 'disp', 'hp', 'wt', 'acc', 'yr']
for col in numeric_cols:
    # pd.to_numeric will convert to a number, coercing errors to NaN
    # errors='coerce' is crucial here: it turns any value that *cannot* be converted
    # into a number into NaN. This handles cases where '?' might not have been replaced,
    # or other non-numeric strings exist.
    data[col] = pd.to_numeric(data[col], errors='coerce')
data = data.apply(lambda x: x.fillna(x.median()), axis = 0)
```

```
In [139... data.head()
```

```
Out[139...      mpg  cyl  disp  hp  wt  acc  yr  car_type  origin_america  origin_asia  origin_europe
0   18.0    8  307.0  130.0 3504  12.0  70         0             True          False          False
1   15.0    8  350.0  165.0 3693  11.5  70         0             True          False          False
2   18.0    8  318.0  150.0 3436  11.0  70         0             True          False          False
3   16.0    8  304.0  150.0 3433  12.0  70         0             True          False          False
4   17.0    8  302.0  140.0 3449  10.5  70         0             True          False          False
```

Aim here is the predict the mileage column given

```
In [142... X = data.drop(['mpg'], axis=1) # independent variable
Y = data[['mpg']] #dependent variable
print(X)
print(Y)
```

```

      cyl  disp    hp  wt   acc  yr  car_type  origin_america  origin_asia \
0      8   307.0  130.0 3504 12.0 70      0           True      False
1      8   350.0  165.0 3693 11.5 70      0           True      False
2      8   318.0  150.0 3436 11.0 70      0           True      False
3      8   304.0  150.0 3433 12.0 70      0           True      False
4      8   302.0  140.0 3449 10.5 70      0           True      False
..    ...    ...    ...    ...    ...    ..    ...    ...
393    4   140.0   86.0 2790 15.6 82      1           True      False
394    4    97.0   52.0 2130 24.6 82      1          False      False
395    4   135.0   84.0 2295 11.6 82      1           True      False
396    4   120.0   79.0 2625 18.6 82      1           True      False
397    4   119.0   82.0 2720 19.4 82      1           True      False

```

```

      origin_europe
0           False
1           False
2           False
3           False
4           False
..          ...
393          False
394           True
395          False
396          False
397          False

```

```
[398 rows x 10 columns]
```

```

      mpg
0      18.0
1      15.0
2      18.0
3      16.0
4      17.0
..    ...
393    27.0
394    44.0
395    32.0
396    28.0
397    31.0

```

```
[398 rows x 1 columns]
```

```
In [144... #Scaling the dataset
```

```

X_s = preprocessing.scale(X)
X_s = pd.DataFrame(X_s, columns = X.columns) #converting scaled data into dataframe

Y_s = preprocessing.scale(Y)
Y_s = pd.DataFrame(Y_s, columns = Y.columns) # ideally train and test data

```

```
In [146... X_train, X_test, Y_train, Y_test = train_test_split(X_s, Y_s, test_size=0.30, random_state=1)
X_train.shape
```

```
Out[146... (278, 10)
```

```
In [148... #Fit simple linear model and find coefficients
```

```

regression_model = LinearRegression()
regression_model.fit(X_train, Y_train)

for idx, col_name in enumerate(X_train.columns):
    print('The coefficient for {} is {}'.format(col_name, regression_model.coef_[0][idx]))

intercept = regression_model.intercept_[0]
print('The intercept is {}'.format(intercept))

```

```

The coefficient for cyl is 0.321022385691611
The coefficient for disp is 0.32483430918483897
The coefficient for hp is -0.22916950059437569
The coefficient for wt is -0.7112101905072298
The coefficient for acc is 0.014713682764191237
The coefficient for yr is 0.3755811949510748
The coefficient for car_type is 0.3814769484233099
The coefficient for origin_america is -0.07472247547584178
The coefficient for origin_asia is 0.044515252035677896
The coefficient for origin_europe is 0.04834854953945386
The intercept is 0.019284116103639764

```

```
In [150... ridge_model = Ridge(alpha =0.3)
ridge_model.fit(X_train, Y_train)

print('Ridge coef.{}'.format(ridge_model.coef_))
```

```

Ridge coef. [[ 0.31649043  0.31320707 -0.22876025 -0.70109447  0.01295851  0.37447352
  0.37725608 -0.07423624  0.04441039  0.04784031]]

```

```
In [155... lasso_model = Lasso(alpha =0.1)
lasso_model.fit(X_train, Y_train)

print('Lasso coef.{}'.format(lasso_model.coef_))
```

```

Lasso coef. [-0.         -0.         -0.01690287 -0.51890013  0.         0.28138241
  0.1278489  -0.01642647  0.         0.         ]

```

```
**Score comparison
```

```
In [160... #Model score - r^2 or coeff of determinant
#r^2 = 1-(RSS/TSS) = Regression error/TSS
```

```

#Simple Linear Model
print(regression_model.score(X_train, Y_train))
print(regression_model.score(X_test, Y_test))

```

```
print('*****')
#Ridge
print(ridge_model.score(X_train, Y_train))
print(ridge_model.score(X_test, Y_test))

print('*****')
#Lasso
print(lasso_model.score(X_train, Y_train))
print(lasso_model.score(X_test, Y_test))
```

```
0.8343770256960538
0.8513421387780066
*****
0.8343617931312616
0.8518882171608506
*****
0.7938010766228453
0.8375229615977083
```

In [170...] *#Model parameter tuning*

```
data_train_test = pd.concat([X_train, Y_train], axis =1)
data_train_test.head()
```

Out[170...]

	cyl	disp	hp	wt	acc	yr	car_type	origin_america	origin_asia	origin_europe	mpg
350	-0.856321	-0.849116	-1.081977	-0.893172	-0.242570	1.351199	0.941412	0.773559	-0.497643	-0.461968	1.432898
59	-0.856321	-0.925936	-1.317736	-0.847061	2.879909	-1.085858	0.941412	-1.292726	-0.497643	2.164651	-0.065919
120	-0.856321	-0.695475	0.201600	-0.121101	-0.024722	-0.815074	0.941412	-1.292726	-0.497643	2.164651	-0.578335
12	1.498191	1.983643	1.197027	0.934732	-2.203196	-1.627426	-1.062235	0.773559	-0.497643	-0.461968	-1.090751
349	-0.856321	-0.983552	-0.951000	-1.165111	0.156817	1.351199	0.941412	-1.292726	2.009471	-0.461968	1.356035

In [174...]

```
import statsmodels.formula.api as smf
ols1 = smf.ols(formula = 'mpg ~ cyl+disp+hp+wt+acc+yr+car_type+origin_america+origin_europe+origin_asia', data = data_train_test).fit()
ols1.params
```

Out[174...]

```
Intercept      0.019284
cyl            0.321022
disp           0.324834
hp            -0.229170
wt            -0.711210
acc           0.014714
yr            0.375581
car_type       0.381477
origin_america -0.074722
origin_europe  0.048349
origin_asia    0.044515
dtype: float64
```

In [179...]

```
print(ols1.summary())
```

```

OLS Regression Results
=====
Dep. Variable:      mpg      R-squared:      0.834
Model:              OLS      Adj. R-squared:    0.829
Method:              Least Squares      F-statistic:    150.0
Date:                Mon, 14 Jul 2025      Prob (F-statistic): 3.12e-99
Time:                21:45:10      Log-Likelihood:   -146.89
No. Observations:    278      AIC:              313.8
Df Residuals:        268      BIC:              350.1
Df Model:             9
Covariance Type:     nonrobust
=====
               coef      std err      t      P>|t|      [0.025      0.975]
-----
Intercept      0.0193      0.025      0.765      0.445      -0.030      0.069
cyl            0.3210      0.112      2.856      0.005      0.100      0.542
disp           0.3248      0.128      2.544      0.012      0.073      0.576
hp            -0.2292      0.079      -2.915      0.004      -0.384      -0.074
wt            -0.7112      0.088      -8.118      0.000      -0.884      -0.539
acc           0.0147      0.039      0.373      0.709      -0.063      0.092
yr            0.3756      0.029     13.088      0.000      0.319      0.432
car_type       0.3815      0.067      5.728      0.000      0.250      0.513
origin_america -0.0747      0.020     -3.723      0.000     -0.114     -0.035
origin_europe  0.0483      0.021      2.270      0.024      0.006      0.090
origin_asia    0.0445      0.020      2.175      0.031      0.004      0.085
=====
Omnibus:          22.678      Durbin-Watson:      2.105
Prob(Omnibus):    0.000      Jarque-Bera (JB):    36.139
Skew:             0.513      Prob(JB):            1.42e-08
Kurtosis:         4.438      Cond. No.            1.59e+16
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The smallest eigenvalue is 6.14e-30. This might indicate that there are
strong multicollinearity problems or that the design matrix is singular.
```

In [184...]

```
# Lets check Sum of Squared Errors (SSE) by predicting value of y for test cases and subtracting from the actual y for the test cases
mse = np.mean((regression_model.predict(X_test)-Y_test)**2)

# root of mean_sq_error is standard deviation i.e. avg variance between predicted and actual
import math
rmse = math.sqrt(mse)
print('Root Mean Squared Error: {}'.format(rmse))
```



Root Mean Squared Error: 0.37766934254087847

In [197...]

```
# Is OLS a good model ? Lets check the residuals for some of these predictor.
```

```
fig = plt.figure(figsize=(8,6))
sn.residplot(x= X_test['hp'], y= Y_test['mpg'], color='blue', lowess=True )

fig = plt.figure(figsize=(8,6))
sn.residplot(x= X_test['acc'], y= Y_test['mpg'], color='green', lowess=True )
plt.show()
```




```
In [228]: y_pred = regression_model.predict(X_test)

min_value = min(Y_test['mpg'].min(), y_pred.min())
max_value = max(Y_test['mpg'].max(), y_pred.max())

plot_range = np.linspace(min_value,max_value, 100)

plt.plot(plot_range, plot_range, color='red', linestyle='--', label='Perfect Prediction (y=x)')
plt.scatter(Y_test['mpg'], y_pred, color='blue')
plt.show()
```



In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: